

Sustainable Software Engineering Workshop

Saif Latifi

24th November 2025

DRIVE-Health, King's College London

Motivation

- Writing software goes beyond *just coding*.
- We want to write code that can be used *beyond the lifetime* of the project.
- Sustainable software is easier to **maintain, extend, and share**.

So how do we achieve this?

Agenda

Today we will broadly cover a range of techniques, including:

- **Version control** with Git and GitHub
- **Testing** and CI/CD
- Building **APIs** with Flask
- **Building and packaging** with Docker

Version Control

What is Version Control?

- It is used to *track changes* to files over time.
- Helpful if you want to *revert* to a previous version.
- Enables *collaboration* among multiple developers by providing tools to combine changes.

Git and GitHub

What is Git?

- A specific flavour of version control.
- Comes installed with most operating systems (expect Windows of course).
- Command line tool with a rich ecosystem.

What is GitHub?

- A web-based platform for hosting Git repositories.
- Facilitates collaboration by easily sharing code with others.
- Provides additional features like issue tracking, pull requests, and project management tools.

Basic Git Concepts: Repository

A **repository** (or *repo*) is a directory that contains your **project files** and the entire history of changes made to those files.

It also contains a hidden subdirectory called `.git` which stores all the *metadata* for your project.

You can create a **new** repository using the command:
`git init`.

Or, **clone** an existing one using `git clone <url>`.

Basic Git Concepts: Commit

A **commit** is a **snapshot** of your repository at a specific point in time. It allows you to see the **changes** made since the last commit, known as the *diff*.

A commit contains a **message** that describes the changes made, along with metadata such as a unique ID, author, and timestamp.

There are a few steps to creating a commit:

1. “Stage” (i.e. select) the changes using `git add <path>`.
2. Create the commit using
`git commit -m "Your commit message".`

Basic Git Concepts: Branch

A **branch** is a **separate** version of your files within the same repository.

It allows you to make changes *without affecting* the main codebase (usually called **main** or **master**).


You can switch branches using
`git switch <branch-name>` and create new branches using the same command with the `-c` flag.

Branch Visualisation

main branch



feature branch



Basic Git Concepts: Branch

A **branch** is a **separate** version of your files within the same repository.

It allows you to make changes *without affecting* the main codebase (usually called **main** or **master**).

You can switch branches using
`git switch <branch-name>` and create new branches using the same command with the `-c` flag.

Basic Git Concepts: Remote

A **remote** is a version of your repository that is **hosted on the internet**.

It allows you to *collaborate* with others by sharing your code and changes. Typically, you would have a single remote called **origin** that points to your GitHub repository.

You can add a remote using

`git remote add <name> <url>` and push changes to it using `git push <remote> <branch>`.

Visualising Git




Figure 1: Git workflow. Source: <https://imgur.com/oodiCnB>

Exercise: Git Basics

See the `exercise/instructions/git.md` for instructions on how to complete this exercise.

Testing

Why Test?

- Tests catch **bugs** early in the development process.
- In addition to this, they can act as a **specification** for your code.
- They provide **confidence** when making changes or adding new features.

Types of Tests

- **Unit tests:** Test **individual functions** or components in isolation.
- **Integration tests:** Test how different **components work together**.
- **End-to-end tests:** Test the **entire application** from start to finish.

In this workshop, we will focus primarily on **unit tests**.

An Example Unit Test

Here is a simple example of a unit test using the `unittest` framework in Python:

```
import unittest
def add(a, b):
    return a + b

class TestAddFunction(unittest.TestCase):
    def test_add_positive(self):
        self.assertEqual(add(2, 3), 5)

    def test_add_negative(self):
        self.assertEqual(add(-1, 1), 0)

if __name__ == '__main__':
    unittest.main()
```

What Should Tests Cover?

- **Normal cases:** Test the **expected** behavior of your functions.
- **Edge cases:** Test the **boundaries** of your functions (e.g., empty inputs, large inputs).
- **Error cases:** Test how your functions handle **invalid inputs** or exceptions.

Example of a Failing Test

From the previous example, here is a test that is expected to fail:

```
import unittest

def add(a, b):
    return a + b

class TestAddFunction(unittest.TestCase):
    def test_add_fails_when_not_numbers(self):
        self.assertRaises(TypeError, add("a", True)) # FAILS!
```

Fixing the Failing Test

To fix the failing test, we need to modify the `add` function to raise a `TypeError` when the inputs are not numbers:

```
import unittest

def add(a, b):
    if not isinstance(a, int) or not isinstance(b, int):
        raise TypeError("Both arguments must be numbers")
    return a + b

class TestAddFunction(unittest.TestCase):
    def test_add_fails_when_not_numbers(self):
        with self.assertRaises(TypeError):
            add("a", True)
```

Exercise: Writing Unit Tests

See the `exercise/instructions/testing.md` for instructions on how to complete this exercise.

Intro to CI/CD

- **Continuous Integration (CI)**: Automatically running **tests and checks** on your code whenever changes are made.
- **Continuous Deployment (CD)**: Automatically **deploying** your code to a production environment.

GitHub provides a built-in CI/CD service called **GitHub Actions** which we will explore next.

Structure of a GitHub Actions Workflow

A GitHub Actions workflow is defined in a YAML file located in the `.github/workflows` directory of your repository.

Here is a simple example of a workflow that echoes “Hello, World!” whenever code is pushed to the repository:

```
name: Hello World

on: push

jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - name: Say Hello
        run: echo "Hello, World!"
```

Exercise: Writing a CI Workflow

See the `exercise/instructions/ci.md` for instructions on how to complete this exercise.

APIs with Flask

What is an API?

An **API** (Application Programming Interface) is a set of rules and protocols that allows different software applications to **communicate** with each other.

The APIs we discuss today will operate over the HTTP protocol and can be used to **access data or perform operations** on a server.

Advantages of Using APIs for Research

- **Modularity:** APIs allow different parts of a system to be developed and maintained **independently**.
- **Reusability:** APIs can be **reused** across different applications or services.
- **Interoperability:** APIs enable different systems to work together, regardless of their **underlying technologies**.
- **Standardisation:** APIs provide a **consistent interface** for accessing data and services.

Basics of HTTP

HTTP (*Hypertext Transfer Protocol*) is the foundation of data communication on the web. There are a few key concepts to understand:

- **Requests:** Clients send requests to servers to **access resources**. Common HTTP methods include GET, POST, and DELETE.
- **Responses:** Servers respond to requests with **status codes** (e.g., 200 OK, 404 Not Found) and **data** (e.g., JSON, HTML).
- **Endpoints:** Specific **URLs** that represent resources or actions in an API.

What is in a HTTP Request?




Figure 2: Anatomy of a HTTP Request. Source:
<https://developer.mozilla.org/en-US/docs/Web/HTTP/Guides/Messages>

Basics of Flask

Flask is a lightweight web framework for Python that makes it easy to build APIs.

Key features of Flask include:

- Simple routing system (mapping URLs to functions).
- Built-in development server (so you can run the API locally).
- Support for various extensions to add functionality (e.g., database integration, authentication).

A Simple Flask API Example

Here is a simple example of a Flask API that has one endpoint which returns a greeting message:

```
from flask import Flask

app = Flask(__name__)

@app.route('/')
def hello_world():
    return 'Hello, World!'

if __name__ == '__main__':
    app.run(debug=True)
```

URL Paramters

The below example builds on the previous one by adding a URL parameter to the endpoint:

```
from flask import Flask

app = Flask(__name__)

@app.route('/greet/<name>')
def greet(name):
    return f'Hello, {name}!'

if __name__ == '__main__':
    app.run(debug=True)
```

POST Request Example

If we want to **pass data** to the API (more than just URL parameters), we can use a **POST** requests. POST requests typically include a **body** containing the data.

```
from flask import Flask, request

app = Flask(__name__)

@app.route('/square', methods=['POST'])
def square():
    data = request.get_json()
    number = data.get('number', 0)
    return {'result': int(number) ** 2}

if __name__ == '__main__':
    app.run(debug=True)
```

Exercise: Building a Simple Flask API

See the `exercise/instructions/api.md` for instructions on how to complete this exercise.

Docker

What is Docker?

Docker is a platform that allows you to **package** your application and its dependencies into a **container**.

Containers are lightweight, portable, and consistent environments that can run on **any system** with Docker installed.

How Does Docker Work?




Figure 3: Docker in relation to the host system. Source:
<https://www.docker.com/resources/what-container/>

Why Use Docker?

- **Consistency:** Ensures that your application **runs the same way** across different environments (development, testing, production).
- **Isolation:** Containers isolate your application from the host system and other containers, **reducing conflicts**.
- **Portability:** Containers can be easily **shared** and deployed on different systems.
- **Scalability:** Docker makes it easy to scale applications by running **multiple containers**.

So How Do We Build a Docker Container?

To build a Docker container, we need to create a **Dockerfile** that specifies the instructions for building the container image. It is every similar to **running commands** in a terminal.

A Dockerfile typically includes:

- A **base image** (e.g., Python, Ubuntu) with the **FROM** instruction.
- Instructions to **install dependencies** (e.g., using **RUN**).
- Commands to **copy your application code** into the container (using **COPY**).
- The command to **run** your application (using **CMD**).

Docker Commands

Here are some common Docker commands:

- `docker build -t <image-name> .`: Build a Docker image from the **Dockerfile** in the current directory.
- `docker run -p <host-port>:<container-port> <image-name>`: Run a Docker **container** from the specified image, mapping ports.
- `docker ps`: **List** running containers.
- `docker stop <container-id>`: **Stop** a running container.

Visualising Docker




Figure 4: Docker architecture. Source: <https://docs.docker.com/get-started/docker-overview/>

A Simple Dockerfile Example

Here is a simple example of a Dockerfile that creates a container for a Python application:

```
# Base image with node preinstalled
FROM node:22-alpine

# Set working directory where commands will be run
WORKDIR /app

# Copy over files
COPY . .

# Install dependencies
RUN npm install

# Run the application when the container starts
CMD ["node", "index.js"]
```

Exercise: Building a Docker Container

See the `exercise/instructions/docker.md` for instructions on how to complete this exercise.

What Happens When We Have Multiple Containers?

When you have multiple containers, you can use **Docker Compose** to define and manage **multi-container** applications.

Docker Compose uses a YAML file (typically called `docker-compose.yml`) to define the different containers (known as *services*) and their configurations. It can also be used to provide a **persistent and shared** storage between containers using *volumes*.

Example Docker Compose File

Here is a simple example of a Docker Compose file that defines two services: a API service and a database.

```
services:  
  app:  
    image: myapp:latest  
    ports:  
      - "5000:5000"  
    depends_on:  
      - db  
  
  db:  
    image: postgres:13  
    # Env variables go here  
    volumes:  
      - db_data:/var/lib/postgresql/data  
  
volumes:  
  db_data:
```

Docker Compose Commands

Here are some common Docker Compose commands:

- `docker compose up`: Start the services defined in the Docker Compose file.
- `docker compose down`: Stop and remove the services.

Conclusion

Summary

In this workshop, we covered several key concepts for sustainable software engineering:

- Version Control with Git and GitHub
- Testing and CI/CD
- Building APIs with Flask
- Building and packaging with Docker

Remember, sustainable software engineering is an **ongoing process** that requires continuous learning and adaptation.